FREE TOOL + GUIDE

Stop Guessing What Your AI Costs

A drop-in tracker for any TypeScript project that shows you exactly where your AI money goes, before the bill surprises you.

MY SPEND BEFORE	MY SPEND AFTER
~$120/mo	**~$50/mo**
37 crons, 12 agents, growing fast, zero visibility	Same features, but now I see every dollar before it leaves

---

Built by Cassy Garner - @cassy.garner
AI systems for content creators

THE PROBLEM

AI costs don't stay flat. They grow.

Every new cron job, every new agent, every new feature adds cost. Without tracking, you don't find out until the bill shocks you.

What was happening in my system

I went from 8 agents to 12 and 24 crons to 37 in a few weeks. My costs were scaling with every feature I shipped, but I had no visibility into any of it. When I finally looked:

- **Apify web scraping.** Found ways to do most of it for free with the Instagram Graph API and YouTube Data API. I only use Apify now for the things I can't replicate.
- **30+ API calls going through the wrong provider.** I'd set up OpenRouter thinking I'd test different models. But once I committed to Anthropic as my primary, the system kept routing calls through OpenRouter anyway, without even checking if Anthropic was available. I was paying OpenRouter markup on calls that should have gone direct.
- **Sonnet routing every Telegram message.** My agent router was using Claude Sonnet to decide which specialist should handle each message. A simple classification task. Haiku does the same thing for 73% less.
- **Crons running every 30 minutes when nothing changed.** My collab reply checker was burning tokens checking an empty inbox.

> **The real problem wasn't what I was wasting today**
>
> It was that costs were growing with every feature and I had no way to see it. I was about to add a lot more agents and automations. Without the tracker, I would have easily hit $350+ within a couple months and had no idea why. The tracker I'm giving you is exactly what I built to catch this before it scaled.

WHAT'S INCLUDED

The AI Cost Tracker

A set of TypeScript files you drop into any project. No database, no third-party service, no setup complexity.

FILE	WHAT IT DOES
`tracker.ts`	Logs every LLM call to a local JSON file with model, tokens, cost, and a label you choose
`wrapper.ts`	Drop-in Anthropic client that tracks automatically. Replace one import and you're done
`pricing.ts`	Cost matrix for every major model (Anthropic, OpenAI, Google, Deepseek). Update when prices change
`summary.ts`	Daily summary with anomaly detection. Flags anything that's 2x+ its normal spend
`review.ts`	Period-over-period review (daily/weekly/monthly). Compares spend, detects spikes, generates suggestions
`ai-review.ts`	AI-powered review. Sends your data to Haiku for specific, contextual recommendations (~$0.01/review)
`cli-*.ts`	Terminal commands: `npm run summary`, `npm run audit`, `npm run review`, `npm run ai-review`

Works with any provider

The wrapper is Anthropic-native, but the tracker works with any LLM. OpenAI, Google, Deepseek, local models. Just call `trackUsage()` after any API call.

HOW TO SET IT UP

Everything happens in Claude Code

Open your project in Claude Code and paste each prompt below. Claude handles the rest. You don't need to know where files go or how imports work.

Prompt 1: Install the tracker

> I want to add AI cost tracking to this project. Here are the source files: https://github.com/cassygarner/ai-cost-tracker
>
> Add these files to my project, install any dependencies they need, and initialize the tracker in my main entry point with a monthly budget cap of $[your budget]. Make sure the usage log file is in my .gitignore so it doesn't get committed.

Prompt 2: Start tracking your AI calls

> Find every place in my codebase where I'm making AI API calls (Anthropic, OpenAI, Google, or any LLM provider) and add cost tracking to each one using the ai-cost-tracker.
>
> Give each call a descriptive label based on what it does, like "classify-email" or "draft-reply" or "generate-summary". I want to be able to see exactly which features cost what.

> **Labels are how you find the waste**
>
> The more specific your labels, the more useful your reports. "classify-email" tells you something. "default" tells you nothing.

AFTER A DAY OR TWO

See where your money goes

Once you've been running for a bit, use these prompts to get your reports.

Prompt 3: Daily cost summary

> Run the ai-cost-tracker daily summary and show me the results.

This shows you:

- Today's spend vs your 7-day average
- Warning if you're above normal
- Anomaly alerts for any feature that spiked
- Spend breakdown by feature and by model

Prompt 4: Find where you're overspending

> Run the ai-cost-tracker model audit and tell me which of my AI calls could use cheaper models without losing quality.

This scans your usage and finds:

- Features using expensive models for simple tasks
- High-frequency callers that could be batched or reduced
- Estimated dollar savings per month

Prompt 5: Set up budget alerts

> Set up the ai-cost-tracker budget alert so I get a notification when I exceed $[your budget] per month. Send the alert to [Telegram/Slack/email]. Here's my webhook/token: [paste it]

TAKE IT FURTHER

Automated cost reviews

The tracker comes with a built-in review system that compares your current spending to the previous period. You pick the frequency: daily, weekly, or monthly.

What the review tells you

- This period's spend vs the previous period (and the % change)
- Your top features ranked by cost, with how much each one changed
- Which models you're using most
- Spikes: anything that's 2x+ its previous cost
- Plain-English suggestions on what to do about it

Prompt 6: Run a review

> Run the ai-cost-tracker review for [daily/weekly/monthly] and show me the results.

Prompt 7: Automate it on a schedule

This is where it gets powerful. Instead of remembering to check, you set it up once and it tells you automatically.

> Set up a cron job that runs the ai-cost-tracker review automatically. I want a [daily/weekly/monthly] review sent to me on [Telegram/Slack/email]. Here's my webhook/token: [paste it]
>
> Only alert me if something spiked or my spend is trending up. If everything looks normal, just log it.

Prompt 8: AI-powered review

The regular review is just math. This one sends your data to an AI model that reads your actual usage patterns and gives you specific recommendations. Things like "your classify-email feature runs 200 times a day on Sonnet, switch it to Haiku and save $4/month."

> Run the ai-cost-tracker AI review for [daily/weekly/monthly]. It should analyze my usage data and give me specific recommendations on what to change.

Cost: about $0.01-0.05 per review (one Haiku call). The AI review tracks its own cost too.

GET THE MOST OUT OF THIS

How to make it actually save you money

1. Label everything

The tracker is only as useful as your labels. "default" tells you nothing. "classify-inbound-email" tells you exactly where to look when costs spike. Be specific.

2. Check your review before you ship new features

Every new AI feature adds cost. Run a review before and after you ship something new so you can see exactly what it added to your bill.

3. Question every expensive model

When the review flags a feature on Sonnet or GPT-4o, ask: "Does this actually need the expensive model?" Test it on Haiku or 4o-mini. If the output is good enough, switch it and save 70-90%.

4. Watch the frequency, not just the price

A cheap call that runs 200 times a day costs more than an expensive call that runs once. The review flags high-frequency features for exactly this reason.

5. Set a budget and respect it

Pick a monthly number you're comfortable with. Set it as your cap. When you get the alert, don't just bump the cap. Figure out what's eating the budget.

> **The automated reviews take a bit more setup**
>
> The cron workflow and notification integrations have more moving parts than the basic tracker. If you get stuck or want to see how others are setting theirs up, the AI Systems Accelerator community is where that happens.

Want me to set this up for you?

If you'd rather have someone build the full system, cron reports, agents, and all, that's what I do. DM me **"BUILD"** on Instagram @cassy.garner

REFERENCE

Where to start with model selection

The single biggest cost saver. Most people use expensive models for everything. Here's where I'd start, but always test with your actual tasks. A model that works for my classification might not work for yours.

TASK TYPE	START WITH	POTENTIAL SAVINGS
Classifying, routing, tagging	Haiku	up to 73% vs Sonnet
Scoring, ranking, filtering	Haiku	up to 73% vs Sonnet
Parsing, extracting structured data	Haiku	up to 73% vs Sonnet
Simple summarization	Haiku / 4o-mini	up to 73-94%
Drafting content, emails	Sonnet	up to 80% vs Opus
Creative writing, strategy	Sonnet	up to 80% vs Opus
Complex reasoning, code gen	Sonnet or Opus	test both

How to actually decide

Run the same 10 prompts through both models. Compare the outputs. Ask yourself:

- Is the cheap model's output good enough for this task?
- Does the expensive model fail less often (fewer retries = lower real cost)?
- How many times per day does this run? A $0.001 difference at 100 calls/day is $3/month.

The cheapest model isn't always the cheapest option. A model that needs 3 retries costs more than one that gets it right the first time. Let the data decide, not the price tag.

Quick cost reference

MODEL	INPUT (PER 1M TOKENS)	OUTPUT (PER 1M TOKENS)
Claude Haiku 4.5	$0.80	$4.00
GPT-4o-mini	$0.15	$0.60
Gemini 2.0 Flash	$0.10	$0.40

MODEL	INPUT (PER 1M TOKENS)	OUTPUT (PER 1M TOKENS)
Claude Sonnet 4	$3.00	$15.00
GPT-4o / GPT-5.4	$2.50	$10.00
Claude Opus 4	$15.00	$75.00

This is what I run on my own system.

37 crons, 12 agents, all tracked. Every dollar accounted for. I built this because I needed it. Now you have it too.

Want to work with me?

I offer three things: consulting to figure out where AI fits in your business, training to get your team actually using it, and full system builds like the one you just read about where I scope it, build it, and hand it over.

If any of that sounds like what you need, DM me **"BUILD"** on Instagram and we'll figure out if it's a good fit.

@cassy.garner on Instagram

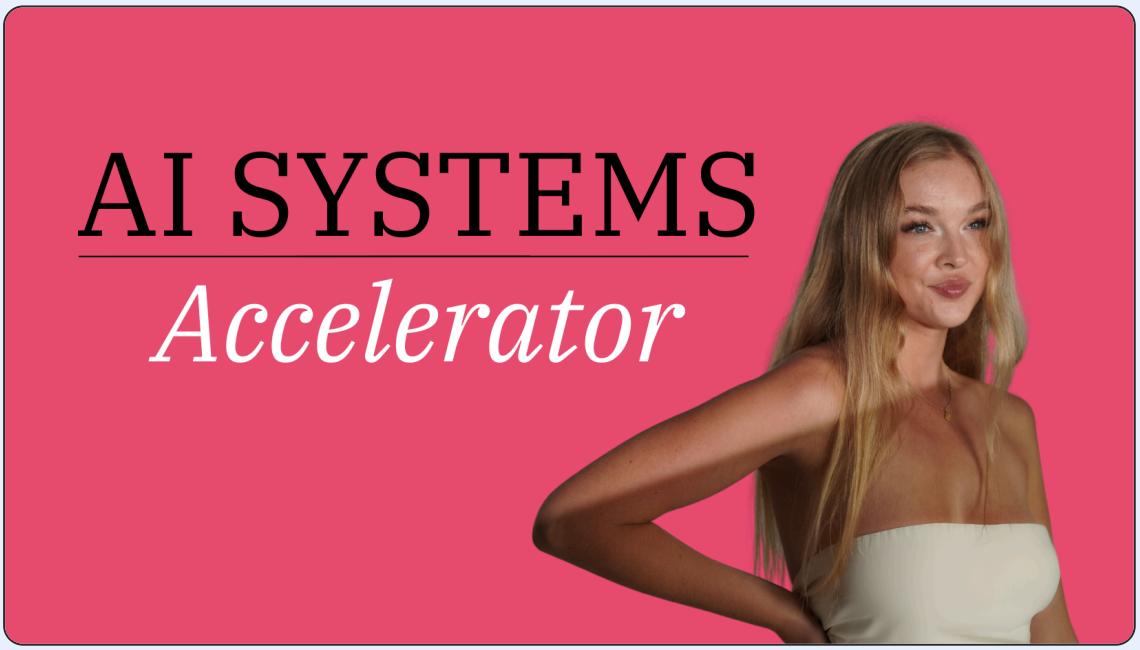

Join the AI Systems Accelerator on Skool for more setups, resources, and a community of likeminded builders.

skool.com/aisystemsaccelerator